Homework of Meachine Learning Techniques: Quiz2

1. Recall that the probabilistic SVM is based on solving the following optimization problem:

$$\min_{A,B} F(A, B) = \frac{1}{N} \sum_{n=1}^{N} \ln\left(1 + \exp\left(-y_n \left(A \cdot (\mathbf{w}_{svm}^T \phi(\mathbf{x}_n) + b_{svm}) + B\right)\right)\right)$$

When using the gradient descent for minimizing $F(A, B)$, we need to compute the gradient first. $z_n = \mathbf{w}_{svm}^T \phi(\mathbf{x}_n) + b_{svm}$, and $p_n = \theta(-y_n(Az_n + B))$, where $\theta(s) = \frac{\exp(s)}{1+\exp(s)}$ is the usual logistic function. What is the gradient $\nabla F(A, B)$?

   ✓ $\frac{1}{N} \sum_{n=1}^{N} [-y_n p_n z_n, -y_n p_n]^T$

   ○ $\frac{1}{N} \sum_{n=1}^{N} [-y_n p_n z_n, +y_n p_n]^T$

   ○ $\frac{1}{N} \sum_{n=1}^{N} [+y_n p_n z_n, -y_n p_n]^T$

   ○ $\frac{1}{N} \sum_{n=1}^{N} [+y_n p_n z_n, +y_n p_n]^T$

   ○ none of the other choices

2. When using the Newton method for minimizing $F(A, B)$ (see Homework 3 of Machine Learning Foundations), we need to compute $-(H(F))^{-1}\nabla F$ in each iteration, where $H(F)$ is the Hessian matrix of F at (A, B). Following the notations of Question 1, what is $H(F)$?

   ○ $\frac{1}{N} \sum_{n=1}^{N} \begin{bmatrix} z_n^2 y_n(1-p_n) & z_n y_n(1-p_n) \\ z_n y_n(1-p_n) & y_n(1-p_n) \end{bmatrix}$

   ○ $\frac{1}{N} \sum_{n=1}^{N} \begin{bmatrix} z_n^2 p_n(1-y_n) & z_n p_n(1-y_n) \\ z_n p_n(1-y_n) & p_n(1-y_n) \end{bmatrix}$

   ○ none of the other choices

   ✓ $\frac{1}{N} \sum_{n=1}^{N} \begin{bmatrix} z_n^2 p_n(1-p_n) & z_n p_n(1-p_n) \\ z_n p_n(1-p_n) & p_n(1-p_n) \end{bmatrix}$

   ○ $\frac{1}{N} \sum_{n=1}^{N} \begin{bmatrix} z_n^2 y_n(1-y_n) & z_n y_n(1-y_n) \\ z_n y_n(1-y_n) & y_n(1-y_n) \end{bmatrix}$

3. Recall that N is the size of the data set and d is the dimensionality of the input space. What is the size of matrix that gets inverted in kernel ridge regression?

   ○ $d \times d$

   ✓ $N \times N$

   ○ $Nd \times Nd$

   ○ $N^2 \times N^2$

   ○ none of the other choices

4. The usual support vector regression model solves the following optimization problem.

$$(P_1) \min_{b,\mathbf{w},\xi^{\vee},\xi^{\wedge}} \frac{1}{2}\mathbf{w}^T\mathbf{w} + C \sum_{n=1}^{N} (\xi_n^{\vee} + \xi_n^{\wedge})$$

$$\text{s.t. } -\xi_n^{\vee} \leq y_n - \mathbf{w}^T\phi(\mathbf{x}_n) - b \leq \epsilon + \xi_n^{\wedge} \xi_n^{\vee} \geq 0, \xi_n^{\wedge} \geq 0.$$

Usual support vector regression penalizes the violations n and n linearly. Another popular formulation, called 2 loss support vector regression in $(P2)$, penalizes the violations quadratically, just like the 2 loss SVM introduced in Homework 1 of Machine Learning Techniques.

$$(P_2) \min_{b,\mathbf{w},\xi^\vee,\xi^\wedge} \frac{1}{2}\mathbf{w}^T\mathbf{w} + C\sum_{n=1}^{N}\left((\xi_n^\vee)^2 + (\xi_n^\wedge)^2\right)$$

$$\text{s.t.} - -\xi_n^\vee \le y_n - \mathbf{w}^T\phi(\mathbf{x}_n) - b \le \epsilon + \xi_n^\wedge.$$

Which of the following is an equivalent 'unconstrained' form of $(P2)$?

○ none of the other choices

○ $\min_{b,\mathbf{w}} \frac{1}{2}\mathbf{w}^T\mathbf{w} + C\sum_{n=1}^{N}(|y_n - \mathbf{w}^T\phi(\mathbf{x}_n) - b| - \epsilon)^2$

√ $\min_{b,\mathbf{w}} \frac{1}{2}\mathbf{w}^T\mathbf{w} + C\sum_{n=1}^{N}(\max(0, |y_n - \mathbf{w}^T\phi(\mathbf{x}_n) - b| - \epsilon))^2$

○ $\min_{b,\mathbf{w}} \frac{1}{2}\mathbf{w}^T\mathbf{w} + C\sum_{n=1}^{N}(\max(\epsilon, |y_n - \mathbf{w}^T\phi(\mathbf{x}_n) - b|))^2$

○ $\min_{b,\mathbf{w}} \frac{1}{2}\mathbf{w}^T\mathbf{w} + C\sum_{n=1}^{N}(y_n - \mathbf{w}^T\phi(\mathbf{x}_n) - b)^2$

5. By a slight modification of the representer theorem presented in the class, the optimal $\mathbf{w}_*$ for $(P2)$ must satisfy $\mathbf{w}_* = \sum_{n=1}^{N}\beta_n\mathbf{z}_n$. We can substitute the form of the optimal $\mathbf{w}_*$ into the answer in Question 4 to derive an optimization problem that contains β (and b) only, which would look like

$$\min_{b,\beta} F(b,\beta) = \frac{1}{2}\sum_{m=1}^{N}\sum_{n=1}^{N}\beta_n\beta_m K(\mathbf{x}_n, \mathbf{x}_m) + \text{something},$$

where $K(\mathbf{x}_n, \mathbf{x}_m) = (\phi(\mathbf{x}_n))^T(\phi(\mathbf{x}_m))$ is the kernel function. One thing that you should see is that $F(b,\beta)$ is differentiable to β_n (and b) and hence you can use gradient descent to solve for the optimal β. For any β, let $s_n = \sum_{m=1}^{N}\beta_m K(\mathbf{x}_n, \mathbf{x}_m) + b$. What is $\frac{\partial F(b,\beta)}{\partial \beta_m}$?

√ $\sum_{n=1}^{N}\beta_n K(\mathbf{x}_n, \mathbf{x}_m) - 2C\sum_{n=1}^{N}[|y_n - s_n| \ge \epsilon](|y_n - s_n| - \epsilon)\mathbf{sign}(y_n - s_n)K(\mathbf{x}_n, \mathbf{x}_m)$

○ $\sum_{n=1}^{N}\beta_n K(\mathbf{x}_n, \mathbf{x}_m) + 2C\sum_{n=1}^{N}[|y_n - s_n| \ge \epsilon](|y_n - s_n| - \epsilon)\text{sign}(y_n - s_n)K(\mathbf{x}_n, \mathbf{x}_m)$

○ $\sum_{n=1}^{N}\beta_n K(\mathbf{x}_n, \mathbf{x}_m) - 2C\sum_{n=1}^{N}[|y_n - s_n| \le \epsilon](|y_n - s_n| - \epsilon)\text{sign}(y_n - s_n)K(\mathbf{x}_n, \mathbf{x}_m)$

○ $\sum_{n=1}^{N}\beta_n K(\mathbf{x}_n, \mathbf{x}_m) + 2C\sum_{n=1}^{N}[|y_n - s_n| \le \epsilon](|y_n - s_n| - \epsilon)\text{sign}(y_n - s_n)K(\mathbf{x}_n, \mathbf{x}_m)$

○ none of the other choices

6. Consider $T+1$ hypotheses $g_0, g_1, \cdots, g_T$. Let $g_0(\mathbf{x}) = 0$ for all x. Assume that your boss holds a test set $\{(\tilde{\mathbf{x}}_m, \tilde{y}_m)\}_{m=1}^{M}$, where you know $\tilde{\mathbf{x}}_m$ but $\tilde{y}_m$ is hidden. Nevertheless, you are allowed to know the squared test error $E_{\text{test}}(g_t) = \frac{1}{M}\sum_{m=1}^{M}(g_t(\tilde{\mathbf{x}}_m) - \tilde{y}_m)^2 = e_t$ for $t = 0, 1, 2, \cdots, T$. Also, assume that $\frac{1}{M}\sum_{m=1}^{M}(g_t(\tilde{\mathbf{x}}_m))^2 = s_t$. Which of the following allows you to calculate $\sum_{m=1}^{M}g_t(\tilde{\mathbf{x}}_m)\tilde{y}_m$? Note that the calculation is the key to the test set blending technique that the NTU team has used in KDDCup2011.

○ $\frac{M}{2}(-e_0 - s_t + e_t)$

○ $\frac{M}{2}(+e_0 - s_t + e_t)$

√ $\frac{M}{2}(+e_0 + s_t - e_t)$

○ none of the other choices

○ $\frac{M}{2}(-e_0 + s_t - e_t)$

7. Consider the case where the target function $f : [0, 1] \to \mathbb{R}$ is given by $f(x) = x^2$ and the input probability distribution is uniform on $[0, 1]$. Assume that the training set has only two examples generated independently from the input probability distribution and noiselessly by f, and the learning model is usual linear regression that minimizes the mean squared error within all hypotheses of the form $h(x) = w_1 x + w_0$. What is $\bar{g}(x)$, the expected value of the hypothesis that the learning algorithm produces (see Page 10 of Lecture 207)?

○ $\bar{g}(x) = 2x - \frac{1}{2}$

○ $\bar{g}(x) = 2x + \frac{1}{2}$

√ $\bar{g}(x) = x - \frac{1}{4}$

○ $\bar{g}(x) = x + \frac{1}{4}$

○ none of the other choices

8. Assume that linear regression (for classification) is used within AdaBoost. That is, we need to solve the weighted-E_{in} optimization problem

$$\min_{\mathbf{w}} E_{in}^{\mathbf{u}}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^{N} u_n (y_n - \mathbf{w}^T \mathbf{x}_n)^2.$$

○ none of the other choices

√ $(\sqrt{u_n} \mathbf{x}_n, \sqrt{u_n} y_n)$

○ $(u_n^{-2} \mathbf{x}_n, u_n^{-2} y_n)$

○ $(u_n^2 \mathbf{x}_n, u_n^2 y_n)$

○ $(u_n \mathbf{x}_n, u_n y_n)$

9. Consider applying the AdaBoost algorithm on a binary classification data set where 99% of the examples are positive. Because there are so many positive examples, the base algorithm within AdaBoost returns a constant classifier $g_1(\mathbf{x}) = +1$ in the first iteration. Let $u_+^{(2)}$ be the individual example weight of each positive example in the second iteration, and $u_-^{(2)}$ be the individual example weight of each negative example in the second iteration. What is $u_+^{(2)}/u_-^{(2)}$?

○ none of the other choices

○ 1/100

√ **1/99**

○ 100

○ 99

10. When talking about non-uniform voting in aggregation, we mentioned that can be viewed as a weight vector learned from any linear algorithm coupled with the following transform:

$$\phi(\mathbf{x}) = \Big(g_1(\mathbf{x}), g_2(\mathbf{x}), \cdots, g_T(\mathbf{x}) \Big).$$

When studying kernel models, we mentioned that the kernel is simply a computational short-cut for the inner product $(\phi(\mathbf{x}))^T (\phi(\mathbf{x}'))$. In this problem, we mix the two topics together using the decision stumps as our $g_t(\mathbf{x})$.

Assume that the input vectors contain only integers between (including) L and R.

$$g_{s,i,\theta}(\mathbf{x}) = s \cdot \text{sign}(x_i - \theta),$$

where $i \in \{1, 2, \cdots, d\}, d$ is the finite dimensionality of the input space,

$$s \in \{-1, +1\}, \theta \in \mathbb{R}, \text{ and sign}(0) = +1$$

Two decision stumps g and $\hat{g}$ are defined as the same if $g(\mathbf{x}) = \hat{g}(\mathbf{x})$ for every $\mathbf{x} \in \mathcal{X}$ Two decision stumps are different if they are not the same. Which of the followings are true?

○ The number of different decision stumps equals the size of $\mathcal{X}$

○ $\mathcal{X}$ is of infinite size

√ **There are 22 different decision stumps for the case of** $d = 2$, $L = 1$, **and** $R = 6$

   $\sqrt{}$ $g_{+1,1,L-1}$ **is the same as** $g_{-1,3,R+1}$

   $\sqrt{}$ $g_{s,i,\theta}$ **is the same as** $g_{s,i,\text{ceiling}(\theta)}$**, where ceiling**$(\theta)$ **is the smallest integer that is greater than or equal to** θ

11. Continuing from the previous question, let $\mathcal{G} = \{$ all different decision stumps for $\mathcal{X}\}$ and enumerate each hypothesis $g \in \mathcal{G}$ by some index t. Define

$$\phi_{ds}(\mathbf{x}) = \Bigg(g_1(\mathbf{x}), g_2(\mathbf{x}), \cdots, g_t(\mathbf{x}), \cdots, g_{|\mathcal{G}|}(\mathbf{x}) \Bigg).$$

Derive a simple equation that evaluates $K_{ds}(\mathbf{x}, \mathbf{x}') = (\phi_{ds}(\mathbf{x}))^T (\phi_{ds}(\mathbf{x}'))$ efficiently. Which of the following equation is correct? Here $\|\mathbf{v}\|_1$ denotes the one-norm of $\mathbf{v}$.

   ◯ $K_{ds}(\mathbf{x}, \mathbf{x}') = 2d(R - L) - 4\|\mathbf{x} - \mathbf{x}'\|_1 - 2$

   ◯ none of the other choices

   $\sqrt{}$ $K_{ds}(\mathbf{x}, \mathbf{x}') = 2d(R - L) - 4\|\mathbf{x} - \mathbf{x}'\|_1 + 2$

   ◯ $K_{ds}(\mathbf{x}, \mathbf{x}') = d(R - L) - 2\|\mathbf{x} - \mathbf{x}'\|_1 - 2$

   ◯ $K_{ds}(\mathbf{x}, \mathbf{x}') = d(R - L) - 2\|\mathbf{x} - \mathbf{x}'\|_1 + 2$

12. For Questions 12-18, implement the AdaBoost-Stump algorithm as introduced in Lecture 208. Run the algorithm on the following set for training: *hw2_adaboost_train.dat* and the following set for testing: *adaboost_test.dat*

    Use a total of $T = 300$ iterations (please do not stop earlier than 300), and calculate Ein and Eout with the 0/1 error.

    For the decision stump algorithm, please implement the following steps. Any ties can be arbitrarily broken.

    1. For any feature i, sort all the $x_{n,i}$ values to $x_{[n],i}$ such that $x_{[n],i} \le x_{[n+1],i}$.

    2. Consider thresholds within $-\infty$ and all the midpoints $\frac{x_{[n],i}+x_{[n+1],i}}{2}$. Test those thresholds with $s \in \{-1, +1\}$ to determine the best (s, θ) combination that minimizes E_{in}^u using feature i.

    3. Pick the best $(s, i,)$ combination by enumerating over all possible i.

    For those interested, Step 2 can be carried out in $O(N)$ time only!!

    Which of the following is true about $E_{in}(g_1)$?

       $\sqrt{}$ $0.2 \le E_{in}(g_1) < 0.3$

       ◯ $E_{in}(g_1) > 0.3$

       ◯ $E_{in}(g_1) = 0$

       ◯ 0

       ◯ $0.1 \le E_{in}(g_1) < 0.2$

13. Which of the following is true about $E_{in}(G)$?

       ◯ $0.1 \le E_{in}(G) < 0.2$

       ◯ $0.2 \le E_{in}(G) < 0.3$

       ◯ $E_{in}(G) > 0.3$

       ◯ $0 < E_{in}(G) < 0.1$

       $\sqrt{}$ $E_{in}(G) = 0$

14. Let $U_t = \sum_{n=1}^{N} u_n^{(t)}$. Which of the following is true about U_2? (note that $U_1 = 1$)

       ◯ $U_2 = 0$

       ◯ $0 < U_2 < 0.1$

       ◯ $0.1 \le U_2 < 0.2$

○ $0.2 \leq U_2 < 0.3$

√ $U_2 > 0.3$

15. Which of the following is true about U_T?

   ○ $U_T = 0$

   √ $0 < U_T < 0.1$

   ○ $0.1 \leq U_T < 0.2$

   ○ $0.2 \leq U_T < 0.3$

   ○ $U_T > 0.3$

16. Which is the following is true about the minimum value of ϵ_t within $t = 1, 2, \cdots, 300$?

   ○ $0 < \text{value} < 0.1$

   ○ $\text{value} > 0.3$

   ○ $\text{value} = 0$

   √ **value** $= 0$

   ○ $0.2 \leq \text{value} < 0.3$

17. Calculate E_{out} with the test set. Which of the following is true about $E_{out}(g_1)$?

   √ $0.2 \leq E_{out}(g_1) < 0.3$

   ○ $E_{out}(g_1) > 0.3$

   ○ $0 < E_{out}(g_1) < 0.1$

   ○ $0.1 \leq E_{out}(g_1) < 0.2$

   ○ $E_{out}(g_1) = 0$

18. Which of the following is true about $E_{out}(G)$?

   √ $0.1 \leq E_{out}(G) < 0.2$

   ○ $0 < E_{out}(G) < 0.1$

   ○ $E_{out}(G) = 0$

   ○ $E_{out}(G) > 0.3$

   ○ $0.2 \leq E_{out}(G) < 0.3$

19. Write a program to implement the kernel ridge regression algorithm from Lecture 206, and use it for classification (i.e. implement LSSVM). Consider the following data set *hw2_lssvm_all.dat* . Use the first 400 examples for training and the remaining for testing. Calculate E_{in} and E_{out} with the 0/1 error. Consider the Gaussian-RBF kernel $\exp\left(-\gamma\|\mathbf{x} - \mathbf{x}'\|^2\right)$ Try all combinations of parameters $\gamma \in \{32, 2, 0.125\}$ and $\lambda \in \{0.001, 1, 1000\}$.
   Among all parameter combinations, which of the following is the range that the minimum $E_{in}(g)$ resides in?

   ○ [0.8,1.0)

   √ **[0,0.2)**

   ○ [0.4,0.6)

   ○ [0.2,0.4)

   ○ [0.6,0.8)

20. Following Question 19, among all parameter combinations, which of the following is the range that the minimum $E_{out}(g)$ resides in?

   √ **[0.2,0.4)**
   ○ [0.8,1.0)
   ○ [0.4,0.6)
   ○ [0.6,0.8)
   ○ [0,0.2)